

## Лабораторная работа №6. Применение сверточных нейронных сетей (многоклассовая классификация)

1. Загрузите данные. Разделите исходный набор данных на обучающую и валидационную выборки.
2. Реализуйте глубокую нейронную сеть со сверточными слоями. Какое качество классификации получено? Какая архитектура сети была использована?
3. Примените дополнение данных (data augmentation). Как это повлияло на качество классификатора?
4. Поэкспериментируйте с готовыми нейронными сетями (например, AlexNet, VGG16, Inception и т.п.), применив передаточное обучение. Как это повлияло на качество классификатора? Можно ли было обойтись без него? Какой максимальный результат удалось получить на контрольной выборке?

```
# Install Kaggle library
!pip install -q kaggle
```

```
import os
os.environ['KAGGLE_USERNAME'] = "awfull1996" # username from the json file
os.environ['KAGGLE_KEY'] = "5e55f76a1cc7a4772bd7803cba8fb2c1" # key from the json file
#!kaggle datasets download -d iarunava/happy-house-dataset # api copied from kaggle
!kaggle datasets download -d datamunge/sign-language-mnist
```

```
↳ Downloading sign-language-mnist.zip to /content
   66% 41.0M/62.6M [00:00<00:00, 26.1MB/s]
  100% 62.6M/62.6M [00:00<00:00, 67.8MB/s]
```

```
!ls
```

```
↳ american_sign_language.PNG  sample_data  sign_mnist_test.csv
   amer_sign2.png             sign-language-mnist.zip  sign_mnist_train
   amer_sign3.png             sign_mnist_test  sign_mnist_train.csv
```

```
!unzip sign-language-mnist.zip
```

```
↳
```

```
Archive:  sign-language-mnist.zip
  inflating: amer_sign2.png
  inflating: amer_sign3.png
  inflating: american_sign_language.PNG
  inflating: sign_mnist_test.csv
  inflating: sign_mnist_test/sign_mnist_test.csv
  inflating: sign_mnist_train.csv
  inflating: sign_mnist_train/sign_mnist_train.csv
```

```
!ls
```

```
↳ american_sign_language.PNG  sample_data  sign_mnist_test.csv
   amer_sign2.png             sign-language-mnist.zip  sign_mnist_train
   amer_sign3.png             sign_mnist_test  sign_mnist_train.csv
```

```
import pandas as pd
df = pd.read_csv('sign_mnist_train.csv')
print(df)

testDF = pd.read_csv('sign_mnist_test.csv')
print(testDF)
```

```
↳
```

	label	pixel11	pixel12	pixel13	...	pixel1781	pixel1782	pixel1783	pixel1784
0	3	107	118	127	...	206	204	203	202
1	6	155	157	156	...	175	103	135	149
2	2	187	188	188	...	198	195	194	195
3	2	211	211	212	...	225	222	229	163
4	13	164	167	170	...	157	163	164	179
...	...	...	...	...	...	...	...	...	...
27450	13	189	189	190	...	234	200	222	225
27451	23	151	154	157	...	195	195	195	194
27452	18	174	174	174	...	203	202	200	200
27453	17	177	181	184	...	47	64	87	93
27454	23	179	180	180	...	197	205	209	215

[27455 rows x 785 columns]

	label	pixel11	pixel12	pixel13	...	pixel1781	pixel1782	pixel1783	pixel1784
0	6	149	149	150	...	106	112	120	107
1	5	126	128	131	...	184	184	182	180
2	10	85	88	92	...	226	225	224	222
3	0	203	205	207	...	230	240	253	255
4	3	188	191	193	...	49	46	46	53
...	...	...	...	...	...	...	...	...	...
7167	1	135	119	108	...	184	176	167	163
7168	12	157	159	161	...	210	210	209	208
7169	2	190	191	190	...	210	211	209	208
7170	4	201	205	208	...	91	67	70	63
7171	2	173	174	173	...	195	195	193	192

[7172 rows x 785 columns]

```
def getData(df):
    y = df["label"].values.reshape((df.shape[0], 1))
    x = df.drop(columns="label").values.reshape((df.shape[0], 28, 28, 1))
    print(y.shape, x.shape)
    return x, y

x, y = getData(df)
tX, tY = getData(testDF)
```



```
(27455, 1) (27455, 28, 28, 1)
(7172, 1) (7172, 28, 28, 1)
```

```
from sklearn.model_selection import train_test_split
x_train, x_val, y_train, y_val = train_test_split(x, y, test_size=.25)
```

**Задание 2.** Реализуйте глубокую нейронную сеть со сверточными слоями. Какое качество классификации получено? Какая архитектура сети была использована?

Была получена точность 93 процента

```
from tensorflow.keras import layers, models, Sequential

model = Sequential([
    layers.Conv2D(64, kernel_size=3, activation='relu', input_shape=(28,28,1)),
    layers.Conv2D(32, kernel_size=5, activation='relu'),
    layers.Dropout(0.25),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu', padding='same'),
    layers.Dropout(0.25),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(128, (3, 3), activation='relu', padding='same'),
    layers.Dropout(0.25),
    layers.MaxPooling2D((2, 2)),
    layers.Flatten(),
    layers.Dense(500, activation='relu'),
    layers.Dropout(0.5),
    layers.Dense(25, activation='softmax')
])

model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])

import tensorflow as tf

model.fit(x_train, y_train, validation_split=0.1, epochs=100, callbacks=[
```

```
tf.keras.callbacks.EarlyStopping(
    patience=2,
    restore_best_weights=True,
    monitor='accuracy'
)
])
```

```
Epoch 1/100
580/580 [=====] - 5s 9ms/step - loss: 3.0719 - accuracy: 0.1263 - val_loss: 1.9795 - val_accu
Epoch 2/100
580/580 [=====] - 5s 8ms/step - loss: 1.1424 - accuracy: 0.6333 - val_loss: 0.3806 - val_accu
Epoch 3/100
580/580 [=====] - 5s 8ms/step - loss: 0.3862 - accuracy: 0.8723 - val_loss: 0.1172 - val_accu
Epoch 4/100
580/580 [=====] - 5s 8ms/step - loss: 0.2489 - accuracy: 0.9207 - val_loss: 0.0915 - val_accu
Epoch 5/100
580/580 [=====] - 5s 8ms/step - loss: 0.1934 - accuracy: 0.9377 - val_loss: 0.0396 - val_accu
Epoch 6/100
580/580 [=====] - 5s 8ms/step - loss: 0.1519 - accuracy: 0.9513 - val_loss: 0.0398 - val_accu
Epoch 7/100
580/580 [=====] - 5s 8ms/step - loss: 0.1315 - accuracy: 0.9608 - val_loss: 0.0433 - val_accu
Epoch 8/100
580/580 [=====] - 5s 8ms/step - loss: 0.1338 - accuracy: 0.9570 - val_loss: 0.0126 - val_accu
Epoch 9/100
580/580 [=====] - 5s 8ms/step - loss: 0.1150 - accuracy: 0.9637 - val_loss: 0.0115 - val_accu
Epoch 10/100
580/580 [=====] - 5s 8ms/step - loss: 0.1006 - accuracy: 0.9700 - val_loss: 0.0176 - val_accu
Epoch 11/100
580/580 [=====] - 5s 8ms/step - loss: 0.1325 - accuracy: 0.9607 - val_loss: 0.0493 - val_accu
Epoch 12/100
580/580 [=====] - 5s 8ms/step - loss: 0.0789 - accuracy: 0.9765 - val_loss: 0.0059 - val_accu
Epoch 13/100
580/580 [=====] - 5s 8ms/step - loss: 0.0865 - accuracy: 0.9752 - val_loss: 0.0184 - val_accu
Epoch 14/100
580/580 [=====] - 5s 8ms/step - loss: 0.0679 - accuracy: 0.9799 - val_loss: 0.0073 - val_accu
Epoch 15/100
580/580 [=====] - 5s 8ms/step - loss: 0.0779 - accuracy: 0.9773 - val_loss: 0.0155 - val_accu
Epoch 16/100
580/580 [=====] - 5s 8ms/step - loss: 0.1046 - accuracy: 0.9718 - val_loss: 0.0025 - val_accu
<tensorflow.python.keras.callbacks.History at 0x7fb1a82672b0>
```

```
model.evaluate(tX, tY)
```

```
↳ 225/225 [=====] - 1s 4ms/step - loss: 0.2069 - accuracy: 0.9334  
[0.20685824751853943, 0.9333519339561462]
```

**Задание 3.** Примените дополнение данных (data augmentation). Как это повлияло на качество классификатора?

После применение качество упало до 5 процентов

```
model = Sequential([
    layers.Conv2D(64, kernel_size=3, activation='relu', input_shape=(28,28,1)),
    layers.Conv2D(32, kernel_size=5, activation='relu'),
    layers.Dropout(0.25),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu', padding='same'),
    layers.Dropout(0.25),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(128, (3, 3), activation='relu', padding='same'),
    layers.Dropout(0.25),
    layers.MaxPooling2D((2, 2)),
    layers.Flatten(),
    layers.Dense(500, activation='relu'),
    layers.Dropout(0.5),
    layers.Dense(25, activation='softmax')
])

model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
```

```
from tensorflow.keras.preprocessing.image import ImageDataGenerator
```

```
datagen = ImageDataGenerator(rotation_range=20, zoom_range=0.15,
    width_shift_range=0.2, height_shift_range=0.2, shear_range=0.15,
    horizontal_flip=True, fill_mode="nearest")
```

```
train_generator = datagen.flow(x_train, y_train, batch_size=32)
```

```
val_generator = datagen.flow(x_val, y_val, batch_size=32)
test_generator = datagen.flow(tX, tY, batch_size=32)
```

```
EPOCHS = 100
```

```
BS = 32
```

```
model.fit_generator(
    train_generator,
    epochs=100,
    validation_data=(x_val,y_val), steps_per_epoch=len(x_train) // 32)
```



```
Epoch 1/100
643/643 [=====] - 13s 20ms/step - loss: 3.2701 - accuracy: 0.0173 - val_loss: 2.9611 - val_ac
Epoch 2/100
643/643 [=====] - 12s 19ms/step - loss: 2.8287 - accuracy: 0.0168 - val_loss: 2.4577 - val_ac
Epoch 3/100
643/643 [=====] - 12s 19ms/step - loss: 2.4743 - accuracy: 0.0523 - val_loss: 1.8554 - val_ac
Epoch 4/100
643/643 [=====] - 12s 19ms/step - loss: 2.0677 - accuracy: 0.0498 - val_loss: 1.4050 - val_ac
Epoch 5/100
643/643 [=====] - 12s 19ms/step - loss: 1.6249 - accuracy: 0.0453 - val_loss: 0.8157 - val_ac
Epoch 6/100
643/643 [=====] - 12s 19ms/step - loss: 1.3345 - accuracy: 0.0450 - val_loss: 0.5770 - val_ac
Epoch 7/100
643/643 [=====] - 12s 19ms/step - loss: 1.1157 - accuracy: 0.0407 - val_loss: 0.4387 - val_ac
Epoch 8/100
643/643 [=====] - 12s 19ms/step - loss: 0.9863 - accuracy: 0.0430 - val_loss: 0.3492 - val_ac
Epoch 9/100
643/643 [=====] - 13s 20ms/step - loss: 0.8957 - accuracy: 0.0443 - val_loss: 0.3064 - val_ac
Epoch 10/100
643/643 [=====] - 12s 19ms/step - loss: 0.8245 - accuracy: 0.0413 - val_loss: 0.2737 - val_ac
Epoch 11/100
643/643 [=====] - 12s 19ms/step - loss: 0.7717 - accuracy: 0.0398 - val_loss: 0.1802 - val_ac
Epoch 12/100
643/643 [=====] - 12s 19ms/step - loss: 0.7431 - accuracy: 0.0396 - val_loss: 0.2054 - val_ac
Epoch 13/100
643/643 [=====] - 13s 20ms/step - loss: 0.6836 - accuracy: 0.0444 - val_loss: 0.2196 - val_ac
Epoch 14/100
643/643 [=====] - 13s 20ms/step - loss: 0.6731 - accuracy: 0.0407 - val_loss: 0.1774 - val_ac
Epoch 15/100
643/643 [=====] - 12s 19ms/step - loss: 0.6326 - accuracy: 0.0411 - val_loss: 0.1585 - val_ac
Epoch 16/100
643/643 [=====] - 12s 19ms/step - loss: 0.6043 - accuracy: 0.0422 - val_loss: 0.1750 - val_ac
Epoch 17/100
643/643 [=====] - 12s 19ms/step - loss: 0.5662 - accuracy: 0.0411 - val_loss: 0.1830 - val_ac
Epoch 18/100
643/643 [=====] - 12s 19ms/step - loss: 0.5747 - accuracy: 0.0390 - val_loss: 0.0813 - val_ac
Epoch 19/100
643/643 [=====] - 12s 19ms/step - loss: 0.5519 - accuracy: 0.0440 - val_loss: 0.0826 - val_ac
Epoch 20/100
643/643 [=====] - 12s 19ms/step - loss: 0.5429 - accuracy: 0.0421 - val_loss: 0.0864 - val_ac
Epoch 21/100
643/643 [=====] - 12s 19ms/step - loss: 0.5425 - accuracy: 0.0417 - val_loss: 0.0843 - val_ac
```



```
Epoch 22/100
643/643 [=====] - 12s 19ms/step - loss: 0.5047 - accuracy: 0.0417 - val_loss: 0.0566 - val_ac
Epoch 23/100
643/643 [=====] - 12s 19ms/step - loss: 0.5330 - accuracy: 0.0411 - val_loss: 0.0739 - val_ac
Epoch 24/100
643/643 [=====] - 12s 19ms/step - loss: 0.5250 - accuracy: 0.0412 - val_loss: 0.0601 - val_ac
Epoch 25/100
643/643 [=====] - 12s 19ms/step - loss: 0.5188 - accuracy: 0.0410 - val_loss: 0.0697 - val_ac
Epoch 26/100
643/643 [=====] - 12s 19ms/step - loss: 0.4970 - accuracy: 0.0429 - val_loss: 0.1668 - val_ac
Epoch 27/100
643/643 [=====] - 12s 19ms/step - loss: 0.5056 - accuracy: 0.0409 - val_loss: 0.0565 - val_ac
Epoch 28/100
643/643 [=====] - 12s 19ms/step - loss: 0.4992 - accuracy: 0.0406 - val_loss: 0.1064 - val_ac
Epoch 29/100
643/643 [=====] - 12s 19ms/step - loss: 0.5036 - accuracy: 0.0407 - val_loss: 0.0907 - val_ac
Epoch 30/100
643/643 [=====] - 12s 19ms/step - loss: 0.4769 - accuracy: 0.0401 - val_loss: 0.0521 - val_ac
Epoch 31/100
643/643 [=====] - 12s 19ms/step - loss: 0.4577 - accuracy: 0.0436 - val_loss: 0.0539 - val_ac
Epoch 32/100
643/643 [=====] - 12s 19ms/step - loss: 0.4644 - accuracy: 0.0401 - val_loss: 0.0487 - val_ac
Epoch 33/100
643/643 [=====] - 13s 19ms/step - loss: 0.4518 - accuracy: 0.0404 - val_loss: 0.0576 - val_ac
Epoch 34/100
643/643 [=====] - 12s 19ms/step - loss: 0.4576 - accuracy: 0.0413 - val_loss: 0.0413 - val_ac
Epoch 35/100
643/643 [=====] - 12s 19ms/step - loss: 0.4462 - accuracy: 0.0415 - val_loss: 0.0423 - val_ac
Epoch 36/100
643/643 [=====] - 12s 19ms/step - loss: 0.4730 - accuracy: 0.0418 - val_loss: 0.0680 - val_ac
Epoch 37/100
643/643 [=====] - 12s 19ms/step - loss: 0.4535 - accuracy: 0.0425 - val_loss: 0.0491 - val_ac
Epoch 38/100
643/643 [=====] - 13s 20ms/step - loss: 0.4598 - accuracy: 0.0407 - val_loss: 0.0609 - val_ac
Epoch 39/100
643/643 [=====] - 13s 19ms/step - loss: 0.4579 - accuracy: 0.0420 - val_loss: 0.0551 - val_ac
Epoch 40/100
643/643 [=====] - 12s 19ms/step - loss: 0.4430 - accuracy: 0.0412 - val_loss: 0.0730 - val_ac
Epoch 41/100
643/643 [=====] - 12s 19ms/step - loss: 0.4553 - accuracy: 0.0429 - val_loss: 0.1224 - val_ac
Epoch 42/100
575/643 [=====>....] - ETA: 1s - loss: 0.4364 - accuracy: 0.0422
```

```
-----
KeyboardInterrupt                                Traceback (most recent call last)
<ipython-input-50-e6d17237f5c8> in <module>()
    13     train_generator,
    14     epochs=100,
--> 15     validation_data=(x_val,y_val), steps_per_epoch=len(x_train) // 32)

----- 10 frames -----
/usr/local/lib/python3.6/dist-packages/tensorflow/python/eager/execute.py in quick_execute(op_name, num_outputs, inputs,
    58     ctx.ensure_initialized()
    59     tensors = pywrap_tfe.TFE_Py_Execute(ctx._handle, device_name, op_name,
--> 60     inputs, attrs, num_outputs)
    61 except core._NotOkStatusException as e:
    62     if name is not None:
```

KeyboardInterrupt:

SEARCH STACK OVERFLOW

```
model.evaluate(tX, tY)
```

```
↳ 225/225 [=====] - 1s 5ms/step - loss: 0.0989 - accuracy: 0.0506
[0.09888710826635361, 0.05061349645256996]
```

**Задание 4.** Поэкспериментируйте с готовыми нейронными сетями (например, AlexNet, VGG16, Inception и т.п.), применив передаточное обучение. Как это повлияло на качество классификатора? Можно ли было обойтись без него? Какой максимальный результат удалось получить на контрольной выборке?

При обучении точность была около 4-х процентов, поэтому дообучивать не было смысла, скорее всего это из-за того что ImageNet не содержит фотографий жестов и не обучался изначально на этих данных, поэтому не стоит использовать передаточное обучение на этом наборе данных

```
from keras.applications import MobileNet
from keras.layers import GlobalAveragePooling2D, Dense
from keras.models import Model
```

```

base_model=MobileNet(weights='imagenet',include_top=False)

x=base_model.output
x=GlobalAveragePooling2D()(x)
x=Dense(1024,activation='relu')(x)
x=Dense(1024,activation='relu')(x)
x=Dense(512,activation='relu')(x)
preds=Dense(25,activation='softmax')(x)
model=Model(inputs=base_model.input,outputs=preds)

print(len(model.layers))
#for layer in model.layers:
#    layer.trainable=False

# or if we want to set the first 20 layers of the network to be non-trainable
for layer in model.layers[:20]:
    layer.trainable=False
for layer in model.layers[20:]:
    layer.trainable=True

```

↳ 92

```

import numpy as np
def transform(dataset):
    newDataset = list()
    for x in dataset:
        x = np.repeat(x, 3, 2)
        newDataset.append(x)
    return np.array(newDataset)

newTrainX = transform(x_train)
newValX = transform(x_val)

print(newTrainX.shape, newValX.shape)

↳ (20591, 28, 28, 3) (6864, 28, 28, 3)

```

```
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
```

```
model.fit(newTrainX, y_train, validation_data=(newValX, y_val), epochs=15)
```

☞ Train on 20591 samples, validate on 6864 samples

```
Epoch 1/15
20591/20591 [=====] - 37s 2ms/step - loss: 3.2080 - accuracy: 0.0454 - val_loss: 3.1988 - va.
Epoch 2/15
20591/20591 [=====] - 32s 2ms/step - loss: 3.1947 - accuracy: 0.0476 - val_loss: 3.1898 - va.
Epoch 3/15
20591/20591 [=====] - 32s 2ms/step - loss: 3.1882 - accuracy: 0.0476 - val_loss: 3.1848 - va.
Epoch 4/15
20591/20591 [=====] - 32s 2ms/step - loss: 3.1845 - accuracy: 0.0476 - val_loss: 3.1819 - va.
Epoch 5/15
20591/20591 [=====] - 32s 2ms/step - loss: 3.1821 - accuracy: 0.0476 - val_loss: 3.1800 - va.
Epoch 6/15
20591/20591 [=====] - 32s 2ms/step - loss: 3.1805 - accuracy: 0.0476 - val_loss: 3.1786 - va.
Epoch 7/15
20591/20591 [=====] - 32s 2ms/step - loss: 3.1794 - accuracy: 0.0476 - val_loss: 3.1777 - va.
Epoch 8/15
20591/20591 [=====] - 32s 2ms/step - loss: 3.1786 - accuracy: 0.0476 - val_loss: 3.1770 - va.
Epoch 9/15
20591/20591 [=====] - 32s 2ms/step - loss: 3.1780 - accuracy: 0.0476 - val_loss: 3.1765 - va.
Epoch 10/15
20591/20591 [=====] - 32s 2ms/step - loss: 3.1776 - accuracy: 0.0476 - val_loss: 3.1761 - va.
Epoch 11/15
20591/20591 [=====] - 32s 2ms/step - loss: 3.1772 - accuracy: 0.0476 - val_loss: 3.1759 - va.
Epoch 12/15
20591/20591 [=====] - 32s 2ms/step - loss: 3.1770 - accuracy: 0.0476 - val_loss: 3.1756 - va.
Epoch 13/15
20591/20591 [=====] - 32s 2ms/step - loss: 3.1768 - accuracy: 0.0476 - val_loss: 3.1754 - va.
Epoch 14/15
20591/20591 [=====] - 32s 2ms/step - loss: 3.1766 - accuracy: 0.0476 - val_loss: 3.1753 - va.
Epoch 15/15
20591/20591 [=====] - 32s 2ms/step - loss: 3.1765 - accuracy: 0.0476 - val_loss: 3.1752 - va.
<keras.callbacks.callbacks.History at 0x7fb19ec2a978>
```

